

A Compiler Optimization Study

Author: Diego Alonso Martínez de Dios

October 27 2019

1 Abstract

Compiler optimization defines the process of improving the results of the compilation pipeline by identifying language patterns so the resulting binaries are easier for the micro-architecture of the CPU to understand and process.

2 Introduction

Performance is an important task that language engineers must take into account during the design phase of the compiler. This can be achieved through bench-marking to evaluate the performance of the compiler on a specific task.

3 Objective

The study will focus on analyzing the assembly code generated by the GCC compiler with and without the Optimize options enabled. By bench-marking the results of the GNU Compiler Collection's with and without optimization features enabled it is expected to result on drastically different results.

4 Development

The study will be split first into the necessary definitions that are necessary to be taken into accounts for the test's development and result interpretation

4.1 Performance

Performance can be defined as the accomplishment of a given task measured against preset known standards of accuracy, completeness, cost and speed. The task of optimization in a compiler is to improve performance. Performance can be improved in different abstraction levels, from the technology used up to the algorithm level, the later one being where the compiler's optimizations work. Given the context of this compiler study we will focus on the performance that the binaries of the C source code written. Performance here has two goals: Latency and Throughput.

4.2 Latency

Time taken to complete a task. Measured in seconds. Lower latencies imply better performances. In order to measure latency t of a program execution, we must consider:

1. Number of executed instructions I .
2. Average number of Clocks Per Instruction (CPI).
3. Clock period T .
4. Clock frequency f .

The equation can be written as:

$$t = I * CPI * T$$

or alternatively as:

$$t = \frac{I * CPI}{f}$$

4.3 Throughput

Number of tasks completed in a given time. Higher throughputs imply better performance.

4.4 Setup

The system where the tests will be executed is a Macbook Pro late 2015. The specs of the processor are as follow:

2.7 GHz (i5-5257U) dual-core Intel Core i5
Broadwell processor with 3 MB shared L3 cache

So the frequency of the clock will be 2.7 GHz To consult the number of clock cycles per instruction of the Intel x86_64 ISA on the Broadwell microarchitecture I used the breakdown study of Agner Fog from the Technical University of Denmark

4.5 C source code

the source code basis that will be put to the test under the optimization process. The program is just a simple function that takes a big array of integers and returns the sum of all the elements in it. This implementation will consider the array passed to the function as a pointer with one million elements of size int reserved in the heap.

The proc.c:

```
#include <stdio.h>
#include <stdlib.h>

int proc(int a[]) {
    int sum = 0, i;
    for (i=0; i < 1000000; i++)
```

```

        sum += a[i];
    return sum;
}

int main(){
    int sum;
    int* arr = (int*)malloc(sizeof(int)*1000000);
    sum = proc(arr);
    printf("Sum: %d\n",sum);
    return 0;
}

```

4.6 Unoptimized Assembly code of the proc function

```

000000000000068a <proc>:
68a: 55                push    %rbp
68b: 48 89 e5          mov     %rsp,%rbp
68e: 48 89 7d e8       mov     %rdi,-0x18(%rbp)
692: c7 45 f8 00 00 00 00 movl    $0x0,-0x8(%rbp)
699: c7 45 fc 00 00 00 00 movl    $0x0,-0x4(%rbp)
6a0: eb 1d            jmp     6bf <proc+0x35>
6a2: 8b 45 fc         mov     -0x4(%rbp),%eax
6a5: 48 98            cltq
6a7: 48 8d 14 85 00 00 00 lea     0x0(,%rax,4),%rdx
6ae: 00
6af: 48 8b 45 e8       mov     -0x18(%rbp),%rax
6b3: 48 01 d0          add     %rdx,%rax
6b6: 8b 00            mov     (%rax),%eax
6b8: 01 45 f8         add     %eax,-0x8(%rbp)
6bb: 83 45 fc 01       addl    $0x1,-0x4(%rbp)
6bf: 81 7d fc 3f 42 0f 00 cmpl    $0xf423f,-0x4(%rbp)
6c6: 7e da            jle     6a2 <proc+0x18>
6c8: 8b 45 f8         mov     -0x8(%rbp),%eax
6cb: 5d              pop     %rbp
6cc: c3              retq

```

From address 0x6a2 to 0x6c6 defines the body of the loop

4.7 O3 Optimized Assembly code of the proc function

```

00000000000007d0 <proc>:
7d0: 48 89 f8          mov     %rdi,%rax
7d3: 48 c1 e8 02       shr     $0x2,%rax
7d7: 48 f7 d8          neg     %rax
7da: 83 e0 03          and     $0x3,%eax
7dd: 0f 84 c7 00 00 00 je      8aa <proc+0xda>
7e3: 83 f8 01          cmp     $0x1,%eax
7e6: 44 8b 17          mov     (%rdi),%r10d
7e9: 0f 84 ab 00 00 00 je      89a <proc+0xca>

```

| | | | |
|------|----------------------|-----------|----------------------|
| 7ef: | 44 03 57 04 | add | 0x4(%rdi),%r10d |
| 7f3: | 83 f8 03 | cmp | \$0x3,%eax |
| 7f6: | 0f 85 be 00 00 00 | jne | 8ba <proc+0xea> |
| 7fc: | 44 03 57 08 | add | 0x8(%rdi),%r10d |
| 800: | 41 b9 3d 42 0f 00 | mov | \$0xf423d,%r9d |
| 806: | be 03 00 00 00 | mov | \$0x3,%esi |
| 80b: | 41 b8 40 42 0f 00 | mov | \$0xf4240,%r8d |
| 811: | 66 0f ef c0 | pxor | %xmm0,%xmm0 |
| 815: | 41 29 c0 | sub | %eax,%r8d |
| 818: | 89 c0 | mov | %eax,%eax |
| 81a: | 44 89 c1 | mov | %r8d,%ecx |
| 81d: | 48 8d 14 87 | lea | (%rdi,%rax,4),%rdx |
| 821: | 31 c0 | xor | %eax,%eax |
| 823: | c1 e9 02 | shr | \$0x2,%ecx |
| 826: | 66 2e 0f 1f 84 00 00 | nopw | %cs:0x0(%rax,%rax,1) |
| 82d: | 00 00 00 | | |
| 830: | 83 c0 01 | add | \$0x1,%eax |
| 833: | 66 0f fe 02 | paddb | (%rdx),%xmm0 |
| 837: | 48 83 c2 10 | add | \$0x10,%rdx |
| 83b: | 39 c1 | cmp | %eax,%ecx |
| 83d: | 77 f1 | ja | 830 <proc+0x60> |
| 83f: | 66 0f 6f c8 | movdqa | %xmm0,%xmm1 |
| 843: | 44 89 c9 | mov | %r9d,%ecx |
| 846: | 66 0f 73 d9 08 | psrldq | \$0x8,%xmm1 |
| 84b: | 66 0f fe c1 | paddb | %xmm1,%xmm0 |
| 84f: | 66 0f 6f c8 | movdqa | %xmm0,%xmm1 |
| 853: | 66 0f 73 d9 04 | psrldq | \$0x4,%xmm1 |
| 858: | 66 0f fe c1 | paddb | %xmm1,%xmm0 |
| 85c: | 66 0f 7e c0 | movd | %xmm0,%eax |
| 860: | 44 01 d0 | add | %r10d,%eax |
| 863: | 45 89 c2 | mov | %r8d,%r10d |
| 866: | 41 83 e2 fc | and | \$0xffffffff,%r10d |
| 86a: | 44 29 d1 | sub | %r10d,%ecx |
| 86d: | 45 39 d0 | cmp | %r10d,%r8d |
| 870: | 41 8d 14 32 | lea | (%r10,%rsi,1),%edx |
| 874: | 74 22 | je | 898 <proc+0xc8> |
| 876: | 48 63 f2 | movslq | %edx,%rsi |
| 879: | 03 04 b7 | add | (%rdi,%rsi,4),%eax |
| 87c: | 83 f9 01 | cmp | \$0x1,%ecx |
| 87f: | 8d 72 01 | lea | 0x1(%rdx),%esi |
| 882: | 74 14 | je | 898 <proc+0xc8> |
| 884: | 48 63 f6 | movslq | %esi,%rsi |
| 887: | 83 c2 02 | add | \$0x2,%edx |
| 88a: | 03 04 b7 | add | (%rdi,%rsi,4),%eax |
| 88d: | 83 f9 02 | cmp | \$0x2,%ecx |
| 890: | 74 06 | je | 898 <proc+0xc8> |
| 892: | 48 63 d2 | movslq | %edx,%rdx |
| 895: | 03 04 97 | add | (%rdi,%rdx,4),%eax |
| 898: | f3 c3 | repz retq | |
| 89a: | 41 b9 3f 42 0f 00 | mov | \$0xf423f,%r9d |

```

8a0:  be 01 00 00 00      mov     $0x1,%esi
8a5:  e9 61 ff ff ff      jmpq    80b <proc+0x3b>
8aa:  41 b9 40 42 0f 00    mov     $0xf4240,%r9d
8b0:  31 f6               xor     %esi,%esi
8b2:  45 31 d2             xor     %r10d,%r10d
8b5:  e9 51 ff ff ff      jmpq    80b <proc+0x3b>
8ba:  41 b9 3e 42 0f 00    mov     $0xf423e,%r9d
8c0:  be 02 00 00 00      mov     $0x2,%esi
8c5:  e9 41 ff ff ff      jmpq    80b <proc+0x3b>
8ca:  66 0f 1f 44 00 00    nopw    0x0(%rax,%rax,1)

```

From address 0x80b to 0x8c5 defines the body of the loop

5 Results

By analyzing the number of times the instructions were executed on the function proc on both cases the following tables were produced: * Along with this document there will be a spreadsheet with the calculations.

5.1 Unoptimized execution

| Instruction | Times Exec | Avg. CPI |
|-------------|------------|----------|
| push | 1 | 1 |
| mov | 3000003 | 0.5 |
| movl | 2 | 1 |
| jmp | 1 | 2 |
| cltq | 1000000 | 1 |
| lea | 1000000 | 1 |
| add | 2000000 | 0.25 |
| addl | 1000000 | 1 |
| cmpl | 1000000 | 0.5 |
| jle | 1000000 | 1 |
| pop | 1 | 4 |
| retq | 1 | 1 |

Then by using the latency formula:

$$t = \frac{I * CPI}{f}$$

We get that the latency t of unoptimized proc is:

$$t = 0.002037041296$$

5.2 Optimized execution

| Instruction | Times Exec | Avg. CPI |
|-------------|------------|----------|
| add | 8000001 | 0.25 |
| and | 1000001 | 0.5 |
| cmp | 4000002 | 0.5 |
| ja | 1000000 | 1 |
| je | 3000002 | 1 |
| jmpq | 3000000 | 1 |
| jne | 1 | 1 |
| lea | 3000000 | 1 |
| mov | 10000004 | 0.5 |
| movd | 1000000 | 0.5 |
| neg | 1 | 0.5 |
| nopw | 2 | 0.5 |
| padd | 3000000 | 0.5 |
| pxor | 1000000 | 0.33 |
| repz | 1000000 | 2 |
| shr | 2000000 | 0.5 |
| sub | 2000000 | 0.25 |
| xor | 3000000 | 0.5 |

Then by using the latency formula:

$$t = \frac{I * CPI}{f}$$

We get that the latency t of optimized proc is:

$$t = 0.009937040093$$

5.3 Comparing results with bash time command

By using the time command we get the following results:

- Unoptimized proc.c

```
time ./proc
real 0m0.002s
user 0m0.001s
sys 0m0.001s
```

- Optimized proc.c

```
time ./proc-03
real 0m0.006s
user 0m0.006s
sys 0m0.000s
```

6 Conclusion

Surprisingly the results were pretty accurate with the calculations made. Compiler optimization is a process that allows better performance of instruction execution in the CPU.

7 References

References

- [1] Lists of instruction latencies, throughputs and micro-operation breakdowns for Intel, AMD, and VIA CPUs https://www.agner.org/optimize/instruction_tables.pdf
- [2] Options That Control Optimization <https://gcc.gnu.org/onlinedocs/gcc/Optimize-Options.html>